



Name: \_\_\_\_\_ Cohort/Batch: \_\_\_\_\_ Date: \_\_\_\_\_ Score: \_\_\_\_\_

## AIRFLOW PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A MLOps/LLMOps

TOTAL: 10 QUESTIONS

**Q1** What is Airflow and what AI/ML use case is it designed for?

Threat Model

---

---

**Q6** How do you evaluate a model's performance in Airflow?

Observability

---

---

**Q2** What are the core abstractions in Airflow?

Architecture

---

---

**Q7** How do you deploy a model trained with Airflow?

Lifecycle

---

---

**Q3** How does Airflow handle training or inference?

Configuration

---

---

**Q8** How does Airflow handle distributed or multi-GPU training?

Compliance

---

---

**Q4** How does Airflow manage data preprocessing?

Hardening

---

---

**Q9** How do you track experiments and versions in Airflow?

Testing

---

---

**Q5** How does Airflow support GPU/TPU acceleration?

SSO / IdP

---

---

**Q10** What are common pitfalls when using Airflow in production?

Tuning

---

---



# ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

## A1 Airflow is a framework or service

Mitigation

Airflow is a framework or service targeting a specific AI/ML domain: training neural networks, building LLM pipelines, deploying models, or orchestrating agents. Understanding its scope helps you know when to use it vs alternatives.

## A6 Evaluation uses a held-out validation set

Observability

Evaluation uses a held-out validation set and metrics (accuracy, F1, BLEU, perplexity) appropriate to the task. Monitor both training and validation metrics to detect overfitting. Use a test set only at the very end to report final results.

## A2 Airflow organizes work around abstractions like

Primitives

Airflow organizes work around abstractions like models, tensors, pipelines, agents, or embeddings. Learning these core types is the first step to using the framework effectively.

## A7 Export the model to a portable

Automation

Export the model to a portable format (ONNX, TorchScript, SavedModel). Serve it via a model server (TorchServe, TF Serving, Triton) or embed it in an API endpoint. Monitor inference latency and drift in production.

## A3 For training, Airflow defines a model

Hardening

For training, Airflow defines a model architecture, a loss function, and an optimizer. For inference, a trained model processes input data and returns predictions. Batch processing and hardware acceleration are key performance levers.

## A8 Airflow provides data-parallel or model-parallel

Governance

Airflow provides data-parallel or model-parallel strategies to split work across GPUs. Data parallelism replicates the model on each GPU and averages gradients. Model parallelism splits layers across devices for models too large for one GPU.

## A4 Airflow provides data loaders, tokenizers, or

Gotchas

Airflow provides data loaders, tokenizers, or pipeline components that transform raw data into the tensor or embedding format the model expects. Proper preprocessing is critical for model correctness and training speed.

## A9 Use an experiment tracker (MLflow, Weights

Assessment

Use an experiment tracker (MLflow, Weights & Biases, Comet) alongside Airflow to log hyperparameters, metrics, and artifacts. Track the code version and data version for full reproducibility.

## A5 Airflow detects available accelerators and moves

Federation

Airflow detects available accelerators and moves computation to them. Specify the device explicitly (cuda, mps, tpu) to avoid accidentally running expensive operations on CPU. Mixed-precision (fp16/bf16) further reduces memory and increases throughput.

## A10 Common issues include training-serving skew

Optimization

Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.