



Name: \_\_\_\_\_ Cohort/Batch: \_\_\_\_\_ Date: \_\_\_\_\_ Score: \_\_\_\_\_

AMAZON DOCUMENTDB PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Databases

TOTAL: 10 QUESTIONS

**Q1** What type of database is Amazon DocumentDB and what are its core strengths? What are del

---

---

**Q6** How do you monitor and tune slow queries in Amazon DocumentDB? Observability

---

---

**Q2** How does Amazon DocumentDB guarantee consistency and handle transactions? Consistency and

---

---

**Q7** What backup and restore strategies does Amazon DocumentDB support? Lifecycle

---

---

**Q3** What index types does Amazon DocumentDB support and when do you use each? Indexing and

---

---

**Q8** How do you handle schema migrations in Amazon DocumentDB? Compliance

---

---

**Q4** How do you design a schema or data model in Amazon DocumentDB? Hardening

---

---

**Q9** What security features does Amazon DocumentDB provide? Security

---

---

**Q5** How does Amazon DocumentDB scale horizontally? SSO / IdP

---

---

**Q10** What are common anti-patterns when using Amazon DocumentDB in production? Tuning

---

---



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

**A1 Amazon DocumentDB is a relational, document, key-value, graph, or time-series store.** Its strengths such as ACID guarantees, horizontal scale, or flexible schema guide the workloads it is best suited for.

**A2 Amazon DocumentDB provides either full ACID transactions or eventual consistency** depending on its CAP theorem positioning. Transaction support varies from none (simple K/V stores) to serializable isolation (relational DBs).

**A3 Amazon DocumentDB offers B-tree, hash, full-text, spatial, or composite indexes** depending on the engine. Choose based on query patterns: B-tree for range queries, hash for equality lookups, full-text for search.

**A4 Schema design in Amazon DocumentDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value).** Understanding query needs upfront prevents costly migrations later.

**A5 Amazon DocumentDB scales via replication (read replicas) and sharding (partitioning data by key)** for write-heavy ones. Some Amazon DocumentDB implementations handle this automatically; others require manual configuration.

**A6 Amazon DocumentDB provides a slow query log, explain/explain plan, or profiling tools** to surface inefficient queries. Common fixes include adding indexes, rewriting queries, or increasing cache size.

**A7 Amazon DocumentDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby.** Point-in-time recovery requires WAL/oplog archiving on top of backups.

**A8 Schema migrations in Amazon DocumentDB are managed with versioned migration files** applied in order by a migration tool. Migrations should be idempotent and tested in staging before production to avoid downtime.

**A9 Amazon DocumentDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging.** Always restrict user privileges to the minimum needed and disable anonymous access in production.

**A10 Common mistakes include selecting all columns instead of needed fields, missing indexes on frequently queried columns, running migrations without backups, and storing large blobs directly in the database instead of object storage.**