



Name: Cohort/Batch: Date: Score:

AMAZON DYNAMODB PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Databases

TOTAL: 10 QUESTIONS

Q1 What type of database is Amazon DynamoDB and what are its core strengths? Architecture

Q6 How do you monitor and tune slow queries in Amazon DynamoDB? Observability

Q2 How does Amazon DynamoDB guarantee consistency and handle transactions? Architecture

Q7 What backup and restore strategies does Amazon DynamoDB support? Reliability

Q3 What index types does Amazon DynamoDB support and when do you use each? Performance

Q8 How do you handle schema migrations in Amazon DynamoDB? Compliance

Q4 How do you design a schema or data model in Amazon DynamoDB? Hardening

Q9 What security features does Amazon DynamoDB provide? Testing

Q5 How does Amazon DynamoDB scale horizontally? SSO / IdP

Q10 What are common anti-patterns when using Amazon DynamoDB in production? Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Amazon DynamoDB is a relational, document, key-value, or time-series store. Its strengths such as ACID guarantees, horizontal scale, or flexible schema guide the workloads it is best suited for.

A6 Amazon DynamoDB provides a slow query log, explain/explain plan, or profiling tools to surface inefficient queries. Common fixes include adding indexes, rewriting queries, or increasing cache size.

A2 Amazon DynamoDB provides either full ACID transactions or eventual consistency depending on its CAP theorem positioning. Transaction support varies from none (simple K/V stores) to serializable isolation (relational DBs).

A7 Amazon DynamoDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby. Point-in-time recovery requires WAL/oplog archiving on top of backups.

A3 Amazon DynamoDB offers B-tree, hash, full-text, spatial, or composite indexes depending on the engine. Choose based on query patterns: B-tree for range queries, hash for equality lookups, full-text for search.

A8 Schema migrations in Amazon DynamoDB are managed with versioned migration files applied in order by a migration tool. Migrations should be idempotent and tested in staging before production to avoid downtime.

A4 Schema design in Amazon DynamoDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value). Understanding query needs upfront prevents costly migrations later.

A9 Amazon DynamoDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging. Always restrict user privileges to the minimum needed and disable anonymous access in production.

A5 Amazon DynamoDB scales via replication (read replicas) or sharding (partitioning data by key) for read-heavy workloads and sharding (partitioning data by key) for write-heavy ones. Some Amazon DynamoDB implementations handle this automatically; others require manual configuration.

A10 Common mistakes include selecting all columns instead of needed fields, missing indexes on frequently queried columns, running migrations without backups, and storing large blobs directly in the database instead of object storage.