



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

AMAZON EKS PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A DevOps

TOTAL: 10 QUESTIONS

Q1 What problem does Amazon EKS solve in DevOps or Threat Model infrastructure?

Q6 How does Amazon EKS integrate with CI/CD pipelines or observability?

Q2 What is Amazon EKS's core architecture? **Architecture**

Q7 How do you debug a failed Amazon EKS deployment? **Lifecycle**

Q3 How does Amazon EKS define infrastructure or configuration as code?

Q8 What are security best practices for Amazon EKS? **Compliance**

Q4 How does Amazon EKS ensure idempotency? **Hardening**

Q9 How does Amazon EKS scale for large environments? **Testing**

Q5 How do you handle secrets in Amazon EKS? **SSO / IdP**

Q10 What are common mistakes teams make when adopting Amazon EKS? **Tuning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Amazon EKS addresses a specific concern

Mitigation

Amazon EKS addresses a specific concern provisioning, configuration management, orchestration, or CI/CD by automating manual steps that are error-prone, slow, or not reproducible when done by hand.

A6 CI/CD pipelines call Amazon EKS in automated

Observability

CI/CD pipelines call Amazon EKS in automated steps: plan on a pull request to preview changes and apply on merge to main. Approval gates can require a human to review the plan before changes go to production.

A2 Amazon EKS has a control plane that

Primitives

Amazon EKS has a control plane that stores desired state and a data plane (agents or push-based SSH) that enforces it. Understanding this distinction clarifies where configuration is evaluated and where changes are applied.

A7 Check Amazon EKS's logs for the specific

Automation

Check Amazon EKS's logs for the specific error message and the resource that failed. For infrastructure tools, re-run with increased verbosity (-v, --debug). Review the state file or event history to understand what changed before the failure.

A3 Amazon EKS uses declarative files (YAML, HCL,

Hardening

Amazon EKS uses declarative files (YAML, HCL, or a DSL) to express desired state. A plan/diff phase shows what will change before applying. Version-controlling these files enables peer review and rollback.

A8 Rotate credentials used by Amazon EKS, restrict

Governance

Rotate credentials used by Amazon EKS, restrict network access to its control plane, enable audit logging of all operations, and use least-privilege service accounts. Never run Amazon EKS with admin credentials in production.

A4 Amazon EKS operations are designed to produce

Gotchas

Amazon EKS operations are designed to produce the same result when run multiple times. Applying the same config twice converges to the desired state rather than creating duplicate resources. This makes automation safe to re-run.

A9 Large Amazon EKS deployments use workspaces or

Assessment

Large Amazon EKS deployments use workspaces or namespaces to isolate environments, remote state backends for team collaboration, and modular code to avoid a single monolithic config that is slow to plan/apply.

A5 Secrets should never be stored in plain

Federation

Secrets should never be stored in plain text in Amazon EKS config files. Use a secrets backend (Vault, AWS Secrets Manager) and inject values at runtime via environment variables or encrypted vaults supported by the tool.

A10 Common pitfalls include storing secrets in plain

Optimization

Common pitfalls include storing secrets in plain text config, skipping the plan step before apply, not reviewing state drift, and not having a rollback strategy when a deployment fails partway through.