



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

APACHE ICEBERG PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Data Platforms

TOTAL: 10 QUESTIONS

Q1 What is Apache Iceberg and what problem in Data Platforms does it solve?

Q6 How does Apache Iceberg handle reliability and high availability?

Q2 What are the core components or concepts in Apache Iceberg?

Q7 What observability does Apache Iceberg provide out of the box?

Q3 How does Apache Iceberg compare to its main alternatives?

Q8 What are common performance bottlenecks in Apache Iceberg?

Q4 How do you get started with Apache Iceberg for a new project?

Q9 What security best practices apply when running Apache Iceberg?

Q5 What configuration options most affect Apache Iceberg's behavior in production?

Q10 What mistakes do teams commonly make when adopting Apache Iceberg?



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Apache Iceberg is a tool or platform

Mitigation

Apache Iceberg is a tool or platform that automates or simplifies a specific concern in the Data Platforms domain, reducing manual effort and operational complexity for engineering teams.

A6 Apache Iceberg provides HA through leader election

Observability

Apache Iceberg provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A2 Apache Iceberg has key building blocks that

Primitives

Apache Iceberg has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 Apache Iceberg exposes metrics (Prometheus endpoint)

Information

Apache Iceberg exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A3 Apache Iceberg trades off simplicity against feature

Hardening

Apache Iceberg trades off simplicity against feature richness differently than its alternatives. Choose Apache Iceberg when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from misconfigured

Performance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A4 Install Apache Iceberg via its package manager

Gotchas

Install Apache Iceberg via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A9 Run Apache Iceberg with the principle of

Assessment

Run Apache Iceberg with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include concurrency

Configuration

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A10 Common pitfalls include skipping the documentation and

Standardization

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.