WorksheetDownload.com

10 CONCEPTUAL Q&A MESSAGING & STREAMING · PRACTICE WORKSHEET

CONCEPTUAL SERIES

Apache Kafka

Name: ____________ Cohort/Batch: ____________ Date: ____________ Score: ____

APACHE KAFKA PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Messaging & Streaming

TOTAL: 10 QUESTIONS

Q1 What is Apache Kafka and what problem in Messaging & Streaming does it solve? *Threat Model*

Q2 What are the core components or concepts in Apache Kafka? *Architecture*

Q3 How does Apache Kafka compare to its main alternatives? *Configuration*

Q4 How do you get started with Apache Kafka for a new project? *Hardening*

Q5 What configuration options most affect Apache Kafka's behavior in production? *SSO / IdP*

Q6 How does Apache Kafka handle reliability and high availability? *Observability*

Q7 What observability does Apache Kafka provide out of the box? *Lifecycle*

Q8 What are common performance bottlenecks in Apache Kafka? *Compliance*

Q9 What security best practices apply when running Apache Kafka? *Testing*

Q10 What mistakes do teams commonly make when adopting Apache Kafka? *Tuning*

ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Apache Kafka is a tool or platform

Mitigation

Apache Kafka is a tool or platform that automates or simplifies a specific concern in the Messaging & Streaming domain, reducing manual effort and operational complexity for engineering teams.

A2 Apache Kafka has key building blocks that

Primitives

Apache Kafka has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A3 Apache Kafka trades off simplicity against feature

Hardening

Apache Kafka trades off simplicity against feature richness differently than its alternatives. Choose Apache Kafka when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A4 Install Apache Kafka via its package manager

Gotchas

Install Apache Kafka via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A5 Production-critical settings typically include concurrency

Federation

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A6 Apache Kafka provides HA through leader election,

Observability

Apache Kafka provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A7 Apache Kafka exposes metrics (Prometheus endpoint or

Automation

Apache Kafka exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A8 Performance bottlenecks typically stem from misconfigured

Governance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A9 Run Apache Kafka with the principle of

Assessment

Run Apache Kafka with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A10 Common pitfalls include skipping the documentation and

Optimization

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.