



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

APACHE SOLR PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Search

TOTAL: 10 QUESTIONS

Q1 What is Apache Solr and what problem in Search does it solve? **What Model**

Q6 How does Apache Solr handle reliability and high availability? **Availability**

Q2 What are the core components or concepts in Apache Solr? **Architecture**

Q7 What observability does Apache Solr provide out of the box? **Log Cycle**

Q3 How does Apache Solr compare to its main alternatives? **Performance**

Q8 What are common performance bottlenecks in Apache Solr? **Performance**

Q4 How do you get started with Apache Solr for a new project? **Getting Started**

Q9 What security best practices apply when running Apache Solr? **Security**

Q5 What configuration options most affect Apache Solr's behavior in production? **Configuration**

Q10 What mistakes do teams commonly make when adopting Apache Solr? **Tuning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Apache Solr is a tool or platform

Mitigation

Apache Solr is a tool or platform that automates or simplifies a specific concern in the Search domain, reducing manual effort and operational complexity for engineering teams.

A6 Apache Solr provides HA through leader election,

Observability

Apache Solr provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A2 Apache Solr has key building blocks that

Primitives

Apache Solr has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 Apache Solr exposes metrics (Prometheus endpoint or

Automation

Apache Solr exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A3 Apache Solr trades off simplicity against feature

Hardening

Apache Solr trades off simplicity against feature richness differently than its alternatives. Choose Apache Solr when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from misconfigured

Performance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A4 Install Apache Solr via its package manager

Gotchas

Install Apache Solr via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A9 Run Apache Solr with the principle of

Assessment

Run Apache Solr with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include concurrency

Concurrency

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A10 Common pitfalls include skipping the documentation and

Documentation

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.