



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

AZURE DATA LAKE STORAGE PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Data Platforms

TOTAL: 10 QUESTIONS

Q1 What is Azure Data Lake Storage and what problem in Data Model Platforms does it solve?

Q6 How does Azure Data Lake Storage handle reliability and high availability?

Q2 What are the core components or concepts in Azure Data Lake Storage?

Q7 What observability does Azure Data Lake Storage provide out of the box?

Q3 How does Azure Data Lake Storage compare to its main alternatives?

Q8 What are common performance bottlenecks in Azure Data Lake Storage?

Q4 How do you get started with Azure Data Lake Storage for a new project?

Q9 What security best practices apply when running Azure Data Lake Storage?

Q5 What configuration options most affect Azure Data Lake Storage's behavior in production?

Q10 What mistakes do teams commonly make when adopting Azure Data Lake Storage?



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Azure Data Lake Storage is a tool Mitigation Azure Data Lake Storage is a tool or platform that automates or simplifies a specific concern in the Data Platforms domain, reducing manual effort and operational complexity for engineering teams.	A6 Azure Data Lake Storage provides HA through Observability Azure Data Lake Storage provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.
A2 Azure Data Lake Storage has key building Primitives Azure Data Lake Storage has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.	A7 Azure Data Lake Storage exposes metrics (Prometheus) Automation Azure Data Lake Storage exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.
A3 Azure Data Lake Storage trades off simplicity Hardening Azure Data Lake Storage trades off simplicity against feature richness differently than its alternatives. Choose Azure Data Lake Storage when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.	A8 Performance bottlenecks typically stem from misconfigured Performance Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.
A4 Install Azure Data Lake Storage via its Gotchas Install Azure Data Lake Storage via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.	A9 Run Azure Data Lake Storage with the Assessment Run Azure Data Lake Storage with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.
A5 Production-critical settings typically include concurrency Concurrency Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.	A10 Common pitfalls include skipping the documentation and Documentation Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.