



Name: Cohort/Batch: Date: Score:

BENTOML PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A MLOps/LLMOps

TOTAL: 10 QUESTIONS

Q1 What is BentoML and what AI/ML use case is it designed for?

Model

.....
.....

Q6 How do you evaluate a model's performance in BentoML?

Reliability

.....
.....

Q2 What are the core abstractions in BentoML?

Architecture

.....
.....

Q7 How do you deploy a model trained with BentoML?

Lifecycle

.....
.....

Q3 How does BentoML handle training or inference?

Configuration

.....
.....

Q8 How does BentoML handle distributed or multi-GPU training?

Scale

.....
.....

Q4 How does BentoML manage data preprocessing?

Hardening

.....
.....

Q9 How do you track experiments and versions in BentoML?

Testing

.....
.....

Q5 How does BentoML support GPU/TPU acceleration?

SSO / IdP

.....
.....

Q10 What are common pitfalls when using BentoML in production?

Ops

.....
.....



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 BentoML is a framework or service targeting

Mitigation

BentoML is a framework or service targeting a specific AI/ML domain training neural networks, building LLM pipelines, deploying models, or orchestrating agents. Understanding its scope helps you know when to use it vs alternatives.

A6 Evaluation uses a held-out validation set and

Observability

Evaluation uses a held-out validation set and metrics (accuracy, F1, BLEU, perplexity) appropriate to the task. Monitor both training and validation metrics to detect overfitting. Use a test set only at the very end to report final results.

A2 BentoML organizes work around abstractions like models, tensors, pipelines, agents, or embeddings. Learning these core types is the first step to using the framework effectively.

Defensitives

A7 Export the model to a portable format

Automation

Export the model to a portable format (ONNX, TorchScript, SavedModel). Serve it via a model server (TorchServe, TF Serving, Triton) or embed it in an API endpoint. Monitor inference latency and drift in production.

A3 For training, BentoML defines a model architecture,

Hardening

For training, BentoML defines a model architecture, a loss function, and an optimizer. For inference, a trained model processes input data and returns predictions. Batch processing and hardware acceleration are key performance levers.

A8 BentoML provides data-parallel or model-parallel strategies to

Defensives

BentoML provides data-parallel or model-parallel strategies to split work across GPUs. Data parallelism replicates the model on each GPU and averages gradients. Model parallelism splits layers across devices for models too large for one GPU.

A4 BentoML provides data loaders, tokenizers, or pipeline components that transform raw data into the tensor or embedding format the model expects. Proper preprocessing is critical for model correctness and training speed.

Gotchas

A9 Use an experiment tracker (MLflow, Weights & Biases, Comet)

Assessment

Use an experiment tracker (MLflow, Weights & Biases, Comet) alongside BentoML to log hyperparameters, metrics, and artifacts. Track the code version and data version for full reproducibility.

A5 BentoML detects available accelerators and moves computation to them. Specify the device explicitly (cuda, mps, tpu) to avoid accidentally running expensive operations on CPU. Mixed-precision (fp16/bf16) further reduces memory and increases throughput.

Federation

A10 Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.

Optimization

Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.