



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

COCKROACHDB PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Databases

TOTAL: 10 QUESTIONS

Q1 What type of database is CockroachDB and what are its core strengths? Security

Q6 How do you monitor and tune slow queries in CockroachDB? Availability

Q2 How does CockroachDB guarantee consistency and handle transactions? Structure

Q7 What backup and restore strategies does CockroachDB support? Lifecycle

Q3 What index types does CockroachDB support and when do you use each? Integration

Q8 How do you handle schema migrations in CockroachDB? Compliance

Q4 How do you design a schema or data model in CockroachDB? Scaling

Q9 What security features does CockroachDB provide? Testing

Q5 How does CockroachDB scale horizontally? SSO / IdP

Q10 What are common anti-patterns when using CockroachDB in production? Training



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 CockroachDB is a relational, document, key-value, graph, or time-series store.

CockroachDB is a relational, document, key-value, graph, or time-series store. Its strengths such as ACID guarantees, horizontal scale, or flexible schema guide the workloads it is best suited for.

A6 CockroachDB provides a slow query log, explain/explain analyze, or profiling tools to surface inefficient queries.

CockroachDB provides a slow query log, explain/explain plan, or profiling tools to surface inefficient queries. Common fixes include adding indexes, rewriting queries, or increasing cache size.

A2 CockroachDB provides either full ACID transactions or eventual consistency.

CockroachDB provides either full ACID transactions or eventual consistency depending on its CAP theorem positioning. Transaction support varies from none (simple K/V stores) to serializable isolation (relational DBs).

A7 CockroachDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby.

CockroachDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby. Point-in-time recovery requires WAL/oplog archiving on top of backups.

A3 CockroachDB offers B-tree, hash, full-text, spatial, or composite indexes.

CockroachDB offers B-tree, hash, full-text, spatial, or composite indexes depending on the engine. Choose based on query patterns: B-tree for range queries, hash for equality lookups, full-text for search.

A8 Schema migrations in CockroachDB are managed with versioned migration files.

Schema migrations in CockroachDB are managed with versioned migration files applied in order by a migration tool. Migrations should be idempotent and tested in staging before production to avoid downtime.

A4 Schema design in CockroachDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value).

Schema design in CockroachDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value). Understanding query needs upfront prevents costly migrations later.

A9 CockroachDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging.

CockroachDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging. Always restrict user privileges to the minimum needed and disable anonymous access in production.

A5 CockroachDB scales via replication (read replicas) for read-heavy workloads and sharding (partitioning data by key) for write-heavy ones.

CockroachDB scales via replication (read replicas) for read-heavy workloads and sharding (partitioning data by key) for write-heavy ones. Some CockroachDB implementations handle this automatically; others require manual configuration.

A10 Common mistakes include selecting all columns instead of needed fields, missing indexes on frequently queried columns, running migrations without backups, and storing large blobs directly in the database instead of object storage.

Common mistakes include selecting all columns instead of needed fields, missing indexes on frequently queried columns, running migrations without backups, and storing large blobs directly in the database instead of object storage.