



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

COUCHDB PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Databases

TOTAL: 10 QUESTIONS

Q1 What type of database is CouchDB and what are its core strengths? Concept Model

Q6 How do you monitor and tune slow queries in CouchDB? Observability

Q2 How does CouchDB guarantee consistency and handle transactions? Architecture

Q7 What backup and restore strategies does CouchDB support? Deployment

Q3 What index types does CouchDB support and when do you use each? Configuration

Q8 How do you handle schema migrations in CouchDB? Compliance

Q4 How do you design a schema or data model in CouchDB? Hardening

Q9 What security features does CouchDB provide? Testing

Q5 How does CouchDB scale horizontally? SSO / IdP

Q10 What are common anti-patterns when using CouchDB in production? Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 CouchDB is a relational, document, key-value, graph, Mitigation

CouchDB is a relational, document, key-value, graph, or time-series store. Its strengths such as ACID guarantees, horizontal scale, or flexible schema guide the workloads it is best suited for.

A6 CouchDB provides a slow query log, explain/explain plan, or profiling tools to Observability

CouchDB provides a slow query log, explain/explain plan, or profiling tools to surface inefficient queries. Common fixes include adding indexes, rewriting queries, or increasing cache size.

A2 CouchDB provides either full ACID transactions or Primitives

CouchDB provides either full ACID transactions or eventual consistency depending on its CAP theorem positioning. Transaction support varies from none (simple K/V stores) to serializable isolation (relational DBs).

A7 CouchDB supports logical backups (export SQL or Automation

CouchDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby. Point-in-time recovery requires WAL/oplog archiving on top of backups.

A3 CouchDB offers B-tree, hash, full-text, spatial, or Hardening

CouchDB offers B-tree, hash, full-text, spatial, or composite indexes depending on the engine. Choose based on query patterns: B-tree for range queries, hash for equality lookups, full-text for search.

A8 Schema migrations in CouchDB are managed with Governance

Schema migrations in CouchDB are managed with versioned migration files applied in order by a migration tool. Migrations should be idempotent and tested in staging before production to avoid downtime.

A4 Schema design in CouchDB involves normalizing relations Optimization

Schema design in CouchDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value). Understanding query needs upfront prevents costly migrations later.

A9 CouchDB supports role-based access control, encrypted Assessment

CouchDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging. Always restrict user privileges to the minimum needed and disable anonymous access in production.

A5 CouchDB scales via replication (read replicas) for Federation

CouchDB scales via replication (read replicas) for read-heavy workloads and sharding (partitioning data by key) for write-heavy ones. Some CouchDB implementations handle this automatically; others require manual configuration.

A10 Common mistakes include selecting all columns instead Optimization

Common mistakes include selecting all columns instead of needed fields, missing indexes on frequently queried columns, running migrations without backups, and storing large blobs directly in the database instead of object storage.