



Name: Cohort/Batch: Date: Score:

CRI-O PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A DevOps

TOTAL: 10 QUESTIONS

Q1 What problem does CRI-O solve in DevOps or infrastructure? Infrastructure

Q6 How does CRI-O integrate with CI/CD pipelines? Observability

Q2 What is CRI-O's core architecture? Architecture

Q7 How do you debug a failed CRI-O deployment? Lifecycle

Q3 How does CRI-O define infrastructure or configuration as code? Configuration

Q8 What are security best practices for CRI-O? Compliance

Q4 How does CRI-O ensure idempotency? Hardening

Q9 How does CRI-O scale for large environments? Testing

Q5 How do you handle secrets in CRI-O? SSO / IdP

Q10 What are common mistakes teams make when adopting CRI-O? Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 CRI-O addresses a specific concern provisioning,

Mitigation

CRI-O addresses a specific concern provisioning, configuration management, orchestration, or CI/CD by automating manual steps that are error-prone, slow, or not reproducible when done by hand.

A6 CI/CD pipelines call CRI-O in automated steps:

Observability

CI/CD pipelines call CRI-O in automated steps: plan on a pull request to preview changes and apply on merge to main. Approval gates can require a human to review the plan before changes go to production.

A2 CRI-O has a control plane that stores

Primitives

CRI-O has a control plane that stores desired state and a data plane (agents or push-based SSH) that enforces it. Understanding this distinction clarifies where configuration is evaluated and where changes are applied.

A7 Check CRI-O's logs for the specific error

Automation

Check CRI-O's logs for the specific error message and the resource that failed. For infrastructure tools, re-run with increased verbosity (-v, --debug). Review the state file or event history to understand what changed before the failure.

A3 CRI-O uses declarative files (YAML, HCL, or

Hardening

CRI-O uses declarative files (YAML, HCL, or a DSL) to express desired state. A plan/diff phase shows what will change before applying. Version-controlling these files enables peer review and rollback.

A8 Rotate credentials used by CRI-O, restrict network

Governance

Rotate credentials used by CRI-O, restrict network access to its control plane, enable audit logging of all operations, and use least-privilege service accounts. Never run CRI-O with admin credentials in production.

A4 CRI-O operations are designed to produce the

Gotchas

CRI-O operations are designed to produce the same result when run multiple times. Applying the same config twice converges to the desired state rather than creating duplicate resources. This makes automation safe to re-run.

A9 Large CRI-O deployments use workspaces or namespaces

Spaces

Large CRI-O deployments use workspaces or namespaces to isolate environments, remote state backends for team collaboration, and modular code to avoid a single monolithic config that is slow to plan/apply.

A5 Secrets should never be stored in plain

Federation

Secrets should never be stored in plain text in CRI-O config files. Use a secrets backend (Vault, AWS Secrets Manager) and inject values at runtime via environment variables or encrypted vaults supported by the tool.

A10 Common pitfalls include storing secrets in plain

Optimization

Common pitfalls include storing secrets in plain text config, skipping the plan step before apply, not reviewing state drift, and not having a rollback strategy when a deployment fails partway through.