



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

DEFI PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Blockchain/Web3

TOTAL: 10 QUESTIONS

Q1 What is DeFi and what AI/ML use case is it designed for?

Threat Model

Q6 How do you evaluate a model's performance in DeFi?

Observability

Q2 What are the core abstractions in DeFi?

Architecture

Q7 How do you deploy a model trained with DeFi?

Lifecycle

Q3 How does DeFi handle training or inference?

Configuration

Q8 How does DeFi handle distributed or multi-GPU training?

Compliance

Q4 How does DeFi manage data preprocessing?

Hardening

Q9 How do you track experiments and versions in DeFi?

Testing

Q5 How does DeFi support GPU/TPU acceleration?

SSO / IdP

Q10 What are common pitfalls when using DeFi in production?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 DeFi is a framework or service

Mitigation

DeFi is a framework or service targeting a specific AI/ML domain - training neural networks, building LLM pipelines, deploying models, or orchestrating agents. Understanding its scope helps you know when to use it vs alternatives.

A6 Evaluation uses a held-out validation set

Observability

Evaluation uses a held-out validation set and metrics (accuracy, F1, BLEU, perplexity) appropriate to the task. Monitor both training and validation metrics to detect overfitting. Use a test set only at the very end to report final results.

A2 DeFi organizes work around abstractions like

Primitives

DeFi organizes work around abstractions like models, tensors, pipelines, agents, or embeddings. Learning these core types is the first step to using the framework effectively.

A7 Export the model to a portable

Automation

Export the model to a portable format (ONNX, TorchScript, SavedModel). Serve it via a model server (TorchServe, TF Serving, Triton) or embed it in an API endpoint. Monitor inference latency and drift in production.

A3 For training, DeFi defines a model

Hardening

For training, DeFi defines a model architecture, a loss function, and an optimizer. For inference, a trained model processes input data and returns predictions. Batch processing and hardware acceleration are key performance levers.

A8 DeFi provides data-parallel or model-parallel

Governance

DeFi provides data-parallel or model-parallel strategies to split work across GPUs. Data parallelism replicates the model on each GPU and averages gradients. Model parallelism splits layers across devices for models too large for one GPU.

A4 DeFi provides data loaders, tokenizers, or

Gotchas

DeFi provides data loaders, tokenizers, or pipeline components that transform raw data into the tensor or embedding format the model expects. Proper preprocessing is critical for model correctness and training speed.

A9 Use an experiment tracker (MLflow, Weights

Assessment

Use an experiment tracker (MLflow, Weights & Biases, Comet) alongside DeFi to log hyperparameters, metrics, and artifacts. Track the code version and data version for full reproducibility.

A5 DeFi detects available accelerators and moves

Federation

DeFi detects available accelerators and moves computation to them. Specify the device explicitly (cuda, mps, tpu) to avoid accidentally running expensive operations on CPU. Mixed-precision (fp16/bf16) further reduces memory and increases throughput.

A10 Common issues include training-serving skew

Optimization

Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.