



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

DISTRIBUTED CACHING PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Distributed Systems

TOTAL: 10 QUESTIONS

Q1 What is Distributed Caching and what problem in Distributed Systems does it solve?

Threat Model

Q6 How does Distributed Caching handle reliability and high availability?

Observability

Q2 What are the core components or concepts in Distributed Caching?

Architecture

Q7 What observability does Distributed Caching provide out of the box?

Lifecycle

Q3 How does Distributed Caching compare to its main alternatives?

Configuration

Q8 What are common performance bottlenecks in Distributed Caching?

Compliance

Q4 How do you get started with Distributed Caching for a new project?

Hardening

Q9 What security best practices apply when running Distributed Caching?

Testing

Q5 What configuration options most affect Distributed Caching's behavior in production?

SSO / IdP

Q10 What mistakes do teams commonly make when adopting Distributed Caching?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Distributed Caching is a tool or

Mitigation

Distributed Caching is a tool or platform that automates or simplifies a specific concern in the Distributed Systems domain, reducing manual effort and operational complexity for engineering teams.

A6 Distributed Caching provides HA through

Observability

Distributed Caching provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A2 Distributed Caching has key building blocks

Primitives

Distributed Caching has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 Distributed Caching exposes metrics

Automation

Distributed Caching exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A3 Distributed Caching trades off simplicity against

Hardening

Distributed Caching trades off simplicity against feature richness differently than its alternatives. Choose Distributed Caching when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from

Governance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A4 Install Distributed Caching via its package

Gotchas

Install Distributed Caching via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A9 Run Distributed Caching with the principle

Assessment

Run Distributed Caching with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include

Federation

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A10 Common pitfalls include skipping the

Optimization

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.