



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

DYNAMODB PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Cloud

TOTAL: 10 QUESTIONS

Q1 What is DynamoDB and what is its primary purpose in its cloud ecosystem? **Init Model**

Q6 How do you manage configuration and secrets for DynamoDB? **DynamoDB?**

Q2 How does IAM work with DynamoDB? **Architecture**

Q7 What are the main cost drivers for DynamoDB and how do you optimize them? **do you**

Q3 How do you monitor DynamoDB in production? **Configuration**

Q8 How do you implement high availability with DynamoDB? **Appliance**

Q4 How do you scale DynamoDB to handle variable load? **Hardening**

Q9 How does DynamoDB integrate with a CI/CD pipeline? **Testing**

Q5 How do you secure DynamoDB in production? **SSO / IdP**

Q10 What are common pitfalls when first adopting DynamoDB? **uning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 **DynamoDB is a managed cloud service that**

Mitigation

DynamoDB is a managed cloud service that handles a specific infrastructure concern (compute, storage, messaging, AI). Using it shifts operational burden to the cloud provider while you focus on application logic.

A6 **Store sensitive values in a managed secret**

Observability

Store sensitive values in a managed secret store (AWS Secrets Manager, GCP Secret Manager, Azure Key Vault). Inject secrets as environment variables at runtime, never bake them into container images or source code.

A2 **Access to DynamoDB is controlled via IAM**

Primitives

Access to DynamoDB is controlled via IAM roles and policies. Attach the minimum required permissions to a service account or role. Avoid long-lived access keys; use instance roles or Workload Identity Federation instead.

A7 **Cost in DynamoDB typically comes from compute**

Automation

Cost in DynamoDB typically comes from compute hours, data transfer, storage, and request counts. Right-size instances, use reserved capacity for steady-state workloads, and enable lifecycle policies to expire old data.

A3 **DynamoDB emits metrics to the cloud provider's**

Hardening

DynamoDB emits metrics to the cloud provider's monitoring service (CloudWatch, Cloud Monitoring, Azure Monitor). Set alerts on key metrics (error rate, latency, capacity). Use distributed tracing to correlate across services.

A8 **Deploy DynamoDB across multiple availability zones**

Use

Deploy DynamoDB across multiple availability zones. Use managed replicas or active-active configurations. Implement health checks and automatic failover. Test resilience regularly with chaos engineering or failover drills.

A4 **DynamoDB supports auto-scaling policies based on CPU, throughput**

Optimization

DynamoDB supports auto-scaling policies based on CPU, queue depth, or custom metrics. Set minimum and maximum capacity. Horizontal scaling (more instances) is generally preferred over vertical for cloud workloads.

A9 **CI/CD pipelines use the cloud CLI or**

Assessment

CI/CD pipelines use the cloud CLI or SDK to deploy updates to DynamoDB after tests pass. Infrastructure-as-code (Terraform, CDK) defines DynamoDB configuration as version-controlled code deployed via the pipeline.

A5 **Place DynamoDB in a private subnet with**

Federation

Place DynamoDB in a private subnet with no public endpoint where possible. Use VPC security groups or firewall rules to allow only required traffic. Encrypt data at rest and in transit. Rotate credentials regularly.

A10 **Common mistakes include misconfigured IAM policies (too**

Optimization

Common mistakes include misconfigured IAM policies (too permissive), no cost alerting, deploying only in one availability zone, and missing backups. Read the service SLA carefully to understand what the cloud provider guarantees.