



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

DYNATRACE PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Observability

TOTAL: 10 QUESTIONS

Q1 What is Dynatrace and what problem in Observability does it solve?

Threat Model

Q6 How does Dynatrace handle reliability and high availability?

Observability

Q2 What are the core components or concepts in Dynatrace?

Architecture

Q7 What observability does Dynatrace provide out of the box?

Lifecycle

Q3 How does Dynatrace compare to its main alternatives?

Configuration

Q8 What are common performance bottlenecks in Dynatrace?

Compliance

Q4 How do you get started with Dynatrace for a new project?

Hardening

Q9 What security best practices apply when running Dynatrace?

Testing

Q5 What configuration options most affect Dynatrace's behavior in production?

SSO / IdP

Q10 What mistakes do teams commonly make when adopting Dynatrace?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Dynatrace is a tool or platform

Mitigation

Dynatrace is a tool or platform that automates or simplifies a specific concern in the Observability domain, reducing manual effort and operational complexity for engineering teams.

A6 Dynatrace provides HA through leader election,

Observability

Dynatrace provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A2 Dynatrace has key building blocks that

Primitives

Dynatrace has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 Dynatrace exposes metrics (Prometheus

Automation

Dynatrace exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A3 Dynatrace trades off simplicity against feature

Hardening

Dynatrace trades off simplicity against feature richness differently than its alternatives. Choose Dynatrace when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from

Governance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A4 Install Dynatrace via its package manager

Gotchas

Install Dynatrace via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A9 Run Dynatrace with the principle of

Assessment

Run Dynatrace with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include

Federation

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A10 Common pitfalls include skipping the

Optimization

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.