



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

ETHEREUM PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Blockchain/Web3

TOTAL: 10 QUESTIONS

Q1 What is Ethereum and what AI/ML use case is it designed for?

Q6 How do you evaluate a model's performance in Ethereum?

Q2 What are the core abstractions in Ethereum?

Architecture

Q7 How do you deploy a model trained with Ethereum?

Lifecycle

Q3 How does Ethereum handle training or inference?

Configuration

Q8 How does Ethereum handle distributed or multi-GPU training?

Q4 How does Ethereum manage data preprocessing?

Hardening

Q9 How do you track experiments and versions in Ethereum?

Testing

Q5 How does Ethereum support GPU/TPU acceleration?

SSO / IdP

Q10 What are common pitfalls when using Ethereum in production?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Ethereum is a framework or service targeting

Mitigation

Ethereum is a framework or service targeting a specific AI/ML domain training neural networks, building LLM pipelines, deploying models, or orchestrating agents. Understanding its scope helps you know when to use it vs alternatives.

A6 Evaluation uses a held-out validation set and

Observability

Evaluation uses a held-out validation set and metrics (accuracy, F1, BLEU, perplexity) appropriate to the task. Monitor both training and validation metrics to detect overfitting. Use a test set only at the very end to report final results.

A2 Ethereum organizes work around abstractions like models, tensors, pipelines, agents, or embeddings. Learning these core types is the first step to using the framework effectively.

Reliability

A7 Export the model to a portable format

Automation

Export the model to a portable format (ONNX, TorchScript, SavedModel). Serve it via a model server (TorchServe, TF Serving, Triton) or embed it in an API endpoint. Monitor inference latency and drift in production.

A3 For training, Ethereum defines a model architecture,

Hardening

For training, Ethereum defines a model architecture, a loss function, and an optimizer. For inference, a trained model processes input data and returns predictions. Batch processing and hardware acceleration are key performance levers.

A8 Ethereum provides data-parallel or model-parallel strategies

Performance

Ethereum provides data-parallel or model-parallel strategies to split work across GPUs. Data parallelism replicates the model on each GPU and averages gradients. Model parallelism splits layers across devices for models too large for one GPU.

A4 Ethereum provides data loaders, tokenizers, or pipeline components that transform raw data into the tensor or embedding format the model expects. Proper preprocessing is critical for model correctness and training speed.

Gotchas

A9 Use an experiment tracker (MLflow, Weights & Biases, Comet)

Assessment

Use an experiment tracker (MLflow, Weights & Biases, Comet) alongside Ethereum to log hyperparameters, metrics, and artifacts. Track the code version and data version for full reproducibility.

A5 Ethereum detects available accelerators and moves computation to them. Specify the device explicitly (cuda, mps, tpu) to avoid accidentally running expensive operations on CPU. Mixed-precision (fp16/bf16) further reduces memory and increases throughput.

Federation

A10 Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.

Optimization