



Name: Cohort/Batch: Date: Score:

GOOGLE DATAFLOW PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Data Engineering

TOTAL: 10 QUESTIONS

Q1 What is Google Dataflow and what problem in Data Engineering does it solve? **Threat Model**

Q6 How does Google Dataflow handle reliability and high availability? **Observability**

Q2 What are the core components or concepts in Google Dataflow? **Architecture**

Q7 What observability does Google Dataflow provide out of the box? **of the cycle**

Q3 How does Google Dataflow compare to its main alternatives? **Alternatives**

Q8 What are common performance bottlenecks in Google Dataflow? **Compliance**

Q4 How do you get started with Google Dataflow for a new project? **Hardening**

Q9 What security best practices apply when running Google Dataflow? **Testing**

Q5 What configuration options most affect Google Dataflow behavior in production? **SSO / IdP**

Q10 What mistakes do teams commonly make when adopting Google Dataflow? **Tuning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Google Dataflow is a tool or platform

Mitigation

Google Dataflow is a tool or platform that automates or simplifies a specific concern in the Data Engineering domain, reducing manual effort and operational complexity for engineering teams.

A6 Google Dataflow provides HA through leader election

Observability

Google Dataflow provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A2 Google Dataflow has key building blocks that

Primitives

Google Dataflow has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 Google Dataflow exposes metrics (Prometheus endpoint)

Observation

Google Dataflow exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A3 Google Dataflow trades off simplicity against feature

Hardening

Google Dataflow trades off simplicity against feature richness differently than its alternatives. Choose Google Dataflow when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from misconfiguration

Performance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A4 Install Google Dataflow via its package manager

Gotchas

Install Google Dataflow via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A9 Run Google Dataflow with the principle of

Assessment

Run Google Dataflow with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include concurrency

Concurrency

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A10 Common pitfalls include skipping the documentation and

Documentation

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.