



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

INFLUXDB PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Databases

TOTAL: 10 QUESTIONS

Q1 What type of database is InfluxDB and what are its core strengths? Core Model

Q6 How do you monitor and tune slow queries in InfluxDB? Observability

Q2 How does InfluxDB guarantee consistency and handle transactions? Architecture

Q7 What backup and restore strategies does InfluxDB support? Lifecycle

Q3 What index types does InfluxDB support and when do you use each? Indexing

Q8 How do you handle schema migrations in InfluxDB? Compliance

Q4 How do you design a schema or data model in InfluxDB? Hardening

Q9 What security features does InfluxDB provide? Testing

Q5 How does InfluxDB scale horizontally? SSO / IdP

Q10 What are common anti-patterns when using InfluxDB in production? Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 InfluxDB is a relational, document, key-value, graph, or time-series store. Mitigation
InfluxDB is a relational, document, key-value, graph, or time-series store. Its strengths such as ACID guarantees, horizontal scale, or flexible schema guide the workloads it is best suited for.

A6 InfluxDB provides a slow query log, explain/explain plan, or profiling tools to surface inefficient queries. Observability
InfluxDB provides a slow query log, explain/explain plan, or profiling tools to surface inefficient queries. Common fixes include adding indexes, rewriting queries, or increasing cache size.

A2 InfluxDB provides either full ACID transactions or eventual consistency depending on its CAP theorem positioning. Primitives
InfluxDB provides either full ACID transactions or eventual consistency depending on its CAP theorem positioning. Transaction support varies from none (simple K/V stores) to serializable isolation (relational DBs).

A7 InfluxDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby. Automation
InfluxDB supports logical backups (export SQL or BSON), physical backups (filesystem snapshot), and continuous replication to a standby. Point-in-time recovery requires WAL/oplog archiving on top of backups.

A3 InfluxDB offers B-tree, hash, full-text, spatial, or composite indexes depending on the engine. Hardening
InfluxDB offers B-tree, hash, full-text, spatial, or composite indexes depending on the engine. Choose based on query patterns: B-tree for range queries, hash for equality lookups, full-text for search.

A8 Schema migrations in InfluxDB are managed with versioned migration files applied in order by a migration tool. Governance
Schema migrations in InfluxDB are managed with versioned migration files applied in order by a migration tool. Migrations should be idempotent and tested in staging before production to avoid downtime.

A4 Schema design in InfluxDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value). Optimization
Schema design in InfluxDB involves normalizing relations (relational DB) or denormalizing for access patterns (document/key-value). Understanding query needs upfront prevents costly migrations later.

A9 InfluxDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging. Assessment
InfluxDB supports role-based access control, encrypted connections (TLS), encryption at rest, and audit logging. Always restrict user privileges to the minimum needed and disable anonymous access in production.

A5 InfluxDB scales via replication (read replicas) for read-heavy workloads and sharding (partitioning data by key) for write-heavy ones. Federation
InfluxDB scales via replication (read replicas) for read-heavy workloads and sharding (partitioning data by key) for write-heavy ones. Some InfluxDB implementations handle this automatically; others require manual configuration.

A10 Common mistakes include selecting all columns instead of needed fields, missing indexes on frequently queried columns, running migrations without backups, and storing large blobs directly in the database instead of object storage. Optimization