



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

KAFKA PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Distributed Systems

TOTAL: 10 QUESTIONS

Q1 What is Kafka and what problem in Distributed Systems does it solve? [Solved](#)

Q6 How does Kafka handle reliability and high availability? [Observability](#)

Q2 What are the core components or concepts in Kafka? [Architecture](#)

Q7 What observability does Kafka provide out of the box? [Lifecycle](#)

Q3 How does Kafka compare to its main alternatives? [Configuration](#)

Q8 What are common performance bottlenecks in Kafka? [Compliance](#)

Q4 How do you get started with Kafka for a new project? [Hardening](#)

Q9 What security best practices apply when running Kafka? [Testing](#)

Q5 What configuration options most affect Kafka's behavior in production? [SaaS / IDP](#)

Q10 What mistakes do teams commonly make when adopting Kafka? [Tuning](#)



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Kafka is a tool or platform that

Mitigation

Kafka is a tool or platform that automates or simplifies a specific concern in the Distributed Systems domain, reducing manual effort and operational complexity for engineering teams.

A2 Kafka has key building blocks that work

Primitives

Kafka has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A3 Kafka trades off simplicity against feature richness

Hardening

Kafka trades off simplicity against feature richness differently than its alternatives. Choose Kafka when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A4 Install Kafka via its package manager or

Gotchas

Install Kafka via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A5 Production-critical settings typically include concurrency

Concurrency

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A6 Kafka provides HA through leader election, replication, and observability

Kafka provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A7 Kafka exposes metrics (Prometheus endpoint or built-in dashboard)

Automation

Kafka exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A8 Performance bottlenecks typically stem from misconfiguration

Performance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A9 Run Kafka with the principle of least

Assessment

Run Kafka with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A10 Common pitfalls include skipping the documentation and

Documentation

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.