



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

LINKERD PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Kubernetes

TOTAL: 10 QUESTIONS

Q1 What problem does Linkerd solve in DevOps or infrastructure?

Q6 How does Linkerd integrate with CI/CD pipelines?

Q2 What is Linkerd's core architecture?

Q7 How do you debug a failed Linkerd deployment?

Q3 How does Linkerd define infrastructure or configuration as code?

Q8 What are security best practices for Linkerd?

Q4 How does Linkerd ensure idempotency?

Q9 How does Linkerd scale for large environments?

Q5 How do you handle secrets in Linkerd?

Q10 What are common mistakes teams make when adopting Linkerd?



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Linkerd addresses a specific concern provisioning, Mitigation

Linkerd addresses a specific concern provisioning, configuration management, orchestration, or CI/CD by automating manual steps that are error-prone, slow, or not reproducible when done by hand.

A6 CI/CD pipelines call Linkerd in automated steps: Observability

CI/CD pipelines call Linkerd in automated steps: plan on a pull request to preview changes and apply on merge to main. Approval gates can require a human to review the plan before changes go to production.

A2 Linkerd has a control plane that stores Primitives

Linkerd has a control plane that stores desired state and a data plane (agents or push-based SSH) that enforces it. Understanding this distinction clarifies where configuration is evaluated and where changes are applied.

A7 Check Linkerd's logs for the specific error Automation

Check Linkerd's logs for the specific error message and the resource that failed. For infrastructure tools, re-run with increased verbosity (-v, --debug). Review the state file or event history to understand what changed before the failure.

A3 Linkerd uses declarative files (YAML, HCL, or Hardening

Linkerd uses declarative files (YAML, HCL, or a DSL) to express desired state. A plan/diff phase shows what will change before applying. Version-controlling these files enables peer review and rollback.

A8 Rotate credentials used by Linkerd, restrict network Governance

Rotate credentials used by Linkerd, restrict network access to its control plane, enable audit logging of all operations, and use least-privilege service accounts. Never run Linkerd with admin credentials in production.

A4 Linkerd operations are designed to produce the Gotchas

Linkerd operations are designed to produce the same result when run multiple times. Applying the same config twice converges to the desired state rather than creating duplicate resources. This makes automation safe to re-run.

A9 Large Linkerd deployments use workspaces or namespaces Sent

Large Linkerd deployments use workspaces or namespaces to isolate environments, remote state backends for team collaboration, and modular code to avoid a single monolithic config that is slow to plan/apply.

A5 Secrets should never be stored in plain Federation

Secrets should never be stored in plain text in Linkerd config files. Use a secrets backend (Vault, AWS Secrets Manager) and inject values at runtime via environment variables or encrypted vaults supported by the tool.

A10 Common pitfalls include storing secrets in plain Optimization

Common pitfalls include storing secrets in plain text config, skipping the plan step before apply, not reviewing state drift, and not having a rollback strategy when a deployment fails partway through.