



Candidate Name: _____ Interviewer: _____ Date: _____ Score: _____

TOP 10 OPENSEARCH INTERVIEW QUESTIONS

Core mechanisms, architecture, production configs & identity integration

PART 1: QUESTIONS 1 - 5

Q1

What is OpenSearch and what machine learning or AI problems is it designed to solve?

Fundamentals

Answer:

Q2

Explain the core abstractions or APIs in OpenSearch (models, tensors, pipelines, agents).

Architecture

Answer:

Q3

How does OpenSearch handle training: defining a model, specifying a loss, and running optimization?

Configuration

Answer:

Q4

What is OpenSearch's approach to data ingestion, preprocessing, and batching?

Security Ops

Answer:

Q5

How do you evaluate model performance in OpenSearch (metrics, validation sets, cross-validation)?

Integration

Answer:



Candidate Name: _____ Interviewer: _____ Date: _____ Score: _____

TOP 10 OPENSEARCH INTERVIEW QUESTIONS

Credential lifecycle, audit, compliance, vulnerability testing & performance

PART 2: QUESTIONS 6 - 10

Q6 How does OpenSearch support model deployment and serving in a production environment? Key Mgmt

Answer:

Q7 What hardware acceleration does OpenSearch support (GPU, TPU) and how do you configure it? Monitoring

Answer:

Q8 How does OpenSearch handle distributed or multi-GPU training? Compliance

Answer:

Q9 What are common debugging techniques for model training issues in OpenSearch? Pen-Testing

Answer:

Q10 How do you version and experiment-track models in a OpenSearch-based ML workflow? Performance

Answer:
