



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

OPENTELEMETRY PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Kubernetes

TOTAL: 10 QUESTIONS

Q1 What problem does OpenTelemetry solve in DevOps or infrastructure?

Threat Model

Q6 How does OpenTelemetry integrate with CI/CD pipelines?

Observability

Q2 What is OpenTelemetry's core architecture?

Architecture

Q7 How do you debug a failed OpenTelemetry deployment?

Lifecycle

Q3 How does OpenTelemetry define infrastructure or configuration as code?

Configuration

Q8 What are security best practices for OpenTelemetry?

Compliance

Q4 How does OpenTelemetry ensure idempotency?

Hardening

Q9 How does OpenTelemetry scale for large environments?

Testing

Q5 How do you handle secrets in OpenTelemetry?

SSO / IdP

Q10 What are common mistakes teams make when adopting OpenTelemetry?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 OpenTelemetry addresses a specific concern

Mitigation

OpenTelemetry addresses a specific concern provisioning, configuration management, orchestration, or CI/CD by automating manual steps that are error-prone, slow, or not reproducible when done by hand.

A6 CI/CD pipelines call OpenTelemetry in

Observability

CI/CD pipelines call OpenTelemetry in automated steps: plan on a pull request to preview changes and apply on merge to main. Approval gates can require a human to review the plan before changes go to production.

A2 OpenTelemetry has a control plane that

Primitives

OpenTelemetry has a control plane that stores desired state and a data plane (agents or push-based SSH) that enforces it. Understanding this distinction clarifies where configuration is evaluated and where changes are applied.

A7 Check OpenTelemetry's logs for the specific

Automation

Check OpenTelemetry's logs for the specific error message and the resource that failed. For infrastructure tools, re-run with increased verbosity (-v, --debug). Review the state file or event history to understand what changed before the failure.

A3 OpenTelemetry uses declarative files (YAML,

Hardening

OpenTelemetry uses declarative files (YAML, HCL, or a DSL) to express desired state. A plan/diff phase shows what will change before applying. Version-controlling these files enables peer review and rollback.

A8 Rotate credentials used by OpenTelemetry,

Governance

Rotate credentials used by OpenTelemetry, restrict network access to its control plane, enable audit logging of all operations, and use least-privilege service accounts. Never run OpenTelemetry with admin credentials in production.

A4 OpenTelemetry operations are designed to

Gotchas

OpenTelemetry operations are designed to produce the same result when run multiple times. Applying the same config twice converges to the desired state rather than creating duplicate resources. This makes automation safe to re-run.

A9 Large OpenTelemetry deployments use

Assessment

Large OpenTelemetry deployments use workspaces or namespaces to isolate environments, remote state backends for team collaboration, and modular code to avoid a single monolithic config that is slow to plan/apply.

A5 Secrets should never be stored in

Federation

Secrets should never be stored in plain text in OpenTelemetry config files. Use a secrets backend (Vault, AWS Secrets Manager) and inject values at runtime via environment variables or encrypted vaults supported by the tool.

A10 Common pitfalls include storing secrets in

Optimization

Common pitfalls include storing secrets in plain text config, skipping the plan step before apply, not reviewing state drift, and not having a rollback strategy when a deployment fails partway through.