



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

PODMAN PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A DevOps

TOTAL: 10 QUESTIONS

Q1 What problem does Podman solve in DevOps or infrastructure?

Threat Model

Q6 How does Podman integrate with CI/CD pipelines?

Observability

Q2 What is Podman's core architecture?

Architecture

Q7 How do you debug a failed Podman deployment?

Lifecycle

Q3 How does Podman define infrastructure or configuration as code?

Configuration as Code

Q8 What are security best practices for Podman?

Compliance

Q4 How does Podman ensure idempotency?

Hardening

Q9 How does Podman scale for large environments?

Testing

Q5 How do you handle secrets in Podman?

SSO / IdP

Q10 What are common mistakes teams make when adopting Podman?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Podman addresses a specific concern provisioning, Mitigation

Podman addresses a specific concern provisioning, configuration management, orchestration, or CI/CD by automating manual steps that are error-prone, slow, or not reproducible when done by hand.

A6 CI/CD pipelines call Podman in automated steps: Observability

CI/CD pipelines call Podman in automated steps: plan on a pull request to preview changes and apply on merge to main. Approval gates can require a human to review the plan before changes go to production.

A2 Podman has a control plane that stores Primitives

Podman has a control plane that stores desired state and a data plane (agents or push-based SSH) that enforces it. Understanding this distinction clarifies where configuration is evaluated and where changes are applied.

A7 Check Podman's logs for the specific error Automation

Check Podman's logs for the specific error message and the resource that failed. For infrastructure tools, re-run with increased verbosity (-v, --debug). Review the state file or event history to understand what changed before the failure.

A3 Podman uses declarative files (YAML, HCL, or Hardening

Podman uses declarative files (YAML, HCL, or a DSL) to express desired state. A plan/diff phase shows what will change before applying. Version-controlling these files enables peer review and rollback.

A8 Rotate credentials used by Podman, restrict network Governance

Rotate credentials used by Podman, restrict network access to its control plane, enable audit logging of all operations, and use least-privilege service accounts. Never run Podman with admin credentials in production.

A4 Podman operations are designed to produce the Gotchas

Podman operations are designed to produce the same result when run multiple times. Applying the same config twice converges to the desired state rather than creating duplicate resources. This makes automation safe to re-run.

A9 Large Podman deployments use workspaces or namespaces

Large Podman deployments use workspaces or namespaces to isolate environments, remote state backends for team collaboration, and modular code to avoid a single monolithic config that is slow to plan/apply.

A5 Secrets should never be stored in plain Federation

Secrets should never be stored in plain text in Podman config files. Use a secrets backend (Vault, AWS Secrets Manager) and inject values at runtime via environment variables or encrypted vaults supported by the tool.

A10 Common pitfalls include storing secrets in plain Optimization

Common pitfalls include storing secrets in plain text config, skipping the plan step before apply, not reviewing state drift, and not having a rollback strategy when a deployment fails partway through.