



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

PRESTO PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Data Engineering

TOTAL: 10 QUESTIONS

Q1 What is Presto and what problem in Data Engineering does it solve? [Model](#)

Q6 How does Presto handle reliability and high availability? [Availability](#)

Q2 What are the core components or concepts in Presto? [Architecture](#)

Q7 What observability does Presto provide out of the box? [Lifecycle](#)

Q3 How does Presto compare to its main alternatives? [Configuration](#)

Q8 What are common performance bottlenecks in Presto? [Compliance](#)

Q4 How do you get started with Presto for a new project? [Hardening](#)

Q9 What security best practices apply when running Presto? [Testing](#)

Q5 What configuration options most affect Presto's behavior in production? [OS / IDP](#)

Q10 What mistakes do teams commonly make when adopting Presto? [Tuning](#)



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Presto is a tool or platform that

Mitigation

Presto is a tool or platform that automates or simplifies a specific concern in the Data Engineering domain, reducing manual effort and operational complexity for engineering teams.

A6 Presto provides HA through leader election, replication, and observability

Presto provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

A2 Presto has key building blocks that work

Primitives

Presto has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 Presto exposes metrics (Prometheus endpoint or built-in dashboard)

Information

Presto exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

A3 Presto trades off simplicity against feature richness

Hardening

Presto trades off simplicity against feature richness differently than its alternatives. Choose Presto when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from misconfigured connection

Performance

Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

A4 Install Presto via its package manager or

Gotchas

Install Presto via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

A9 Run Presto with the principle of least

Assessment

Run Presto with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include concurrency

Configuration

Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

A10 Common pitfalls include skipping the documentation and

Configuration

Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.