



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

PYSPARK PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Data Engineering

TOTAL: 10 QUESTIONS

Q1 What is PySpark and what problem in Data Engineering does it solve? Introduction

Q6 How does PySpark handle reliability and high availability? Reliability

Q2 What are the core components or concepts in PySpark? Architecture

Q7 What observability does PySpark provide out of the box? Lifecycle

Q3 How does PySpark compare to its main alternatives? Configuration

Q8 What are common performance bottlenecks in PySpark? Compliance

Q4 How do you get started with PySpark for a new project? Hardening

Q9 What security best practices apply when running PySpark? Testing

Q5 What configuration options most affect PySpark's behavior in production? Deployment

Q10 What mistakes do teams commonly make when adopting PySpark? Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 PySpark is a tool or platform that

Mitigation

PySpark is a tool or platform that automates or simplifies a specific concern in the Data Engineering domain, reducing manual effort and operational complexity for engineering teams.

A6 PySpark provides HA through leader election, replication, or stateless horizontal scaling. Deploy at least two instances across availability zones and configure health checks for automatic failover.

Availability

A2 PySpark has key building blocks that work

Primitives

PySpark has key building blocks that work together: a configuration layer defines intent, a runtime layer enforces it, and an API or CLI layer allows operators to manage and observe the system.

A7 PySpark exposes metrics (Prometheus endpoint or built-in dashboard), structured logs, and optionally distributed traces. Define SLOs on the key metrics and alert before users are affected.

Monitoring

A3 PySpark trades off simplicity against feature richness

Hardening

PySpark trades off simplicity against feature richness differently than its alternatives. Choose PySpark when its specific strengths performance, ecosystem, or ease of use align with your team's primary constraints.

A8 Performance bottlenecks typically stem from misconfigured connection pools, large payload sizes, insufficient worker threads, or missing caches. Profile with built-in diagnostics before optimizing.

Performance

A4 Install PySpark via its package manager or binary release. Follow the quickstart to initialize a project, configure the minimal required settings, and run the default command to verify the setup works end-to-end.

Gotchas

A9 Run PySpark with the principle of least

Assessment

Run PySpark with the principle of least privilege, enable TLS for all communication, use a secrets manager for credentials, and keep the software patched to the latest stable release.

A5 Production-critical settings typically include concurrency limits, timeout values, retry policies, resource limits, and logging levels. Review the official production hardening guide before deploying.

Configuration

A10 Common pitfalls include skipping the documentation and learning the API by trial and error, using default configurations in production, lacking a runbook for operational issues, and not testing failover scenarios.

Documentation