



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

RANCHER PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Kubernetes

TOTAL: 10 QUESTIONS

Q1 What problem does Rancher solve in DevOps or infrastructure?

Threat Model

Q6 How does Rancher integrate with CI/CD pipelines?

Observability

Q2 What is Rancher's core architecture?

Architecture

Q7 How do you debug a failed Rancher deployment?

Lifecycle

Q3 How does Rancher define infrastructure or configuration as code?

Configuration

Q8 What are security best practices for Rancher?

Compliance

Q4 How does Rancher ensure idempotency?

Hardening

Q9 How does Rancher scale for large environments?

Testing

Q5 How do you handle secrets in Rancher?

SSO / IdP

Q10 What are common mistakes teams make when adopting Rancher?

Tuning



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Rancher addresses a specific concern

Mitigation

Rancher addresses a specific concern provisioning, configuration management, orchestration, or CI/CD by automating manual steps that are error-prone, slow, or not reproducible when done by hand.

A6 CI/CD pipelines call Rancher in automated

Observability

CI/CD pipelines call Rancher in automated steps: plan on a pull request to preview changes and apply on merge to main. Approval gates can require a human to review the plan before changes go to production.

A2 Rancher has a control plane that

Primitives

Rancher has a control plane that stores desired state and a data plane (agents or push-based SSH) that enforces it. Understanding this distinction clarifies where configuration is evaluated and where changes are applied.

A7 Check Rancher's logs for the specific

Automation

Check Rancher's logs for the specific error message and the resource that failed. For infrastructure tools, re-run with increased verbosity (-v, --debug). Review the state file or event history to understand what changed before the failure.

A3 Rancher uses declarative files (YAML, HCL,

Hardening

Rancher uses declarative files (YAML, HCL, or a DSL) to express desired state. A plan/diff phase shows what will change before applying. Version-controlling these files enables peer review and rollback.

A8 Rotate credentials used by Rancher, restrict

Governance

Rotate credentials used by Rancher, restrict network access to its control plane, enable audit logging of all operations, and use least-privilege service accounts. Never run Rancher with admin credentials in production.

A4 Rancher operations are designed to produce

Gotchas

Rancher operations are designed to produce the same result when run multiple times. Applying the same config twice converges to the desired state rather than creating duplicate resources. This makes automation safe to re-run.

A9 Large Rancher deployments use workspaces or

Assessment

Large Rancher deployments use workspaces or namespaces to isolate environments, remote state backends for team collaboration, and modular code to avoid a single monolithic config that is slow to plan/apply.

A5 Secrets should never be stored in

Federation

Secrets should never be stored in plain-text in Rancher config files. Use a secrets backend (Vault, AWS Secrets Manager) and inject values at runtime via environment variables or encrypted vaults supported by the tool.

A10 Common pitfalls include storing secrets in

Optimization

Common pitfalls include storing secrets in plain-text config, skipping the plan step before apply, not reviewing state drift, and not having a rollback strategy when a deployment fails partway through.