



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

SPEECH AI PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A AI/ML

TOTAL: 10 QUESTIONS

Q1 What is Speech AI and what AI/ML use case is it designed for? **Threat Model**

Q6 How do you evaluate a model's performance in Speech AI? **Observability**

Q2 What are the core abstractions in Speech AI? **Architecture**

Q7 How do you deploy a model trained with Speech AI? **Lifecycle**

Q3 How does Speech AI handle training or inference? **Configuration**

Q8 How does Speech AI handle distributed or multi-GPU training? **Compliance**

Q4 How does Speech AI manage data preprocessing? **Hardening**

Q9 How do you track experiments and versions in Speech AI? **Testing**

Q5 How does Speech AI support GPU/TPU acceleration? **SSO / IdP**

Q10 What are common pitfalls when using Speech AI in production? **Tuning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Speech AI is a framework or

Mitigation

Speech AI is a framework or service targeting a specific AI/ML domain training neural networks, building LLM pipelines, deploying models, or orchestrating agents. Understanding its scope helps you know when to use it vs alternatives.

A6 Evaluation uses a held-out validation set

Observability

Evaluation uses a held-out validation set and metrics (accuracy, F1, BLEU, perplexity) appropriate to the task. Monitor both training and validation metrics to detect overfitting. Use a test set only at the very end to report final results.

A2 Speech AI organizes work around abstractions

Primitives

Speech AI organizes work around abstractions like models, tensors, pipelines, agents, or embeddings. Learning these core types is the first step to using the framework effectively.

A7 Export the model to a portable

Automation

Export the model to a portable format (ONNX, TorchScript, SavedModel). Serve it via a model server (TorchServe, TF Serving, Triton) or embed it in an API endpoint. Monitor inference latency and drift in production.

A3 For training, Speech AI defines a

Hardening

For training, Speech AI defines a model architecture, a loss function, and an optimizer. For inference, a trained model processes input data and returns predictions. Batch processing and hardware acceleration are key performance levers.

A8 Speech AI provides data-parallel or

Governance

Speech AI provides data-parallel or model-parallel strategies to split work across GPUs. Data parallelism replicates the model on each GPU and averages gradients. Model parallelism splits layers across devices for models too large for one GPU.

A4 Speech AI provides data loaders, tokenizers,

Gotchas

Speech AI provides data loaders, tokenizers, or pipeline components that transform raw data into the tensor or embedding format the model expects. Proper preprocessing is critical for model correctness and training speed.

A9 Use an experiment tracker (MLflow, Weights

Assessment

Use an experiment tracker (MLflow, Weights & Biases, Comet) alongside Speech AI to log hyperparameters, metrics, and artifacts. Track the code version and data version for full reproducibility.

A5 Speech AI detects available accelerators and

Federation

Speech AI detects available accelerators and moves computation to them. Specify the device explicitly (cuda, mps, tpu) to avoid accidentally running expensive operations on CPU. Mixed-precision (fp16/bf16) further reduces memory and increases throughput.

A10 Common issues include training-serving skew

Optimization

Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.