



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

TOKEN MONITORING PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A MLOps/LLMOps

TOTAL: 10 QUESTIONS

Q1 What is Token Monitoring and what AI/ML use case is it designed for? **Great Model**

Q6 How do you evaluate a model's performance in Token Monitoring? **Observability**

Q2 What are the core abstractions in Token Monitoring? **Architecture**

Q7 How do you deploy a model trained with Token Monitoring? **Lifecycle**

Q3 How does Token Monitoring handle training or inference? **Duration**

Q8 How does Token Monitoring handle distributed or multi-GPU training? **Scale**

Q4 How does Token Monitoring manage data preprocessing? **Ordering**

Q9 How do you track experiments and versions in Token Monitoring? **Testing**

Q5 How does Token Monitoring support GPU/TPU acceleration? **IdP**

Q10 What are common pitfalls when using Token Monitoring in production? **Tuning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Token Monitoring is a framework or service

Mitigation

Token Monitoring is a framework or service targeting a specific AI/ML domain training neural networks, building LLM pipelines, deploying models, or orchestrating agents. Understanding its scope helps you know when to use it vs alternatives.

A6 Evaluation uses a held-out validation set and

Observability

Evaluation uses a held-out validation set and metrics (accuracy, F1, BLEU, perplexity) appropriate to the task. Monitor both training and validation metrics to detect overfitting. Use a test set only at the very end to report final results.

A2 Token Monitoring organizes work around abstractions like

Alternatives

Token Monitoring organizes work around abstractions like models, tensors, pipelines, agents, or embeddings. Learning these core types is the first step to using the framework effectively.

A7 Export the model to a portable format

Automation

Export the model to a portable format (ONNX, TorchScript, SavedModel). Serve it via a model server (TorchServe, TF Serving, Triton) or embed it in an API endpoint. Monitor inference latency and drift in production.

A3 For training, Token Monitoring defines a model

Hardening

For training, Token Monitoring defines a model architecture, a loss function, and an optimizer. For inference, a trained model processes input data and returns predictions. Batch processing and hardware acceleration are key performance levers.

A8 Token Monitoring provides data-parallel or model-parallel

Performance

Token Monitoring provides data-parallel or model-parallel strategies to split work across GPUs. Data parallelism replicates the model on each GPU and averages gradients. Model parallelism splits layers across devices for models too large for one GPU.

A4 Token Monitoring provides data loaders, tokenizers, or

Gotchas

Token Monitoring provides data loaders, tokenizers, or pipeline components that transform raw data into the tensor or embedding format the model expects. Proper preprocessing is critical for model correctness and training speed.

A9 Use an experiment tracker (MLflow, Weights &

Assessment

Use an experiment tracker (MLflow, Weights & Biases, Comet) alongside Token Monitoring to log hyperparameters, metrics, and artifacts. Track the code version and data version for full reproducibility.

A5 Token Monitoring detects available accelerators and moves

Optimization

Token Monitoring detects available accelerators and moves computation to them. Specify the device explicitly (cuda, mps, tpu) to avoid accidentally running expensive operations on CPU. Mixed-precision (fp16/bf16) further reduces memory and increases throughput.

A10 Common issues include training-serving skew (different

Optimization

Common issues include training-serving skew (different preprocessing), missing input validation, no latency SLA enforcement, models not versioned, and no process to retrain on drifted data distributions.