



Name: _____ Cohort/Batch: _____ Date: _____ Score: _____

VERTEX AI PRACTICE SHEET - 10 CONCEPTUAL Q&A

10 conceptual Q&A Cloud

TOTAL: 10 QUESTIONS

Q1 What is Vertex AI and what is its primary purpose in its cloud ecosystem? **Cloud**

Q6 How do you manage configuration and secrets for Vertex AI? **Security**

Q2 How does IAM work with Vertex AI? **Architecture**

Q7 What are the main cost drivers for Vertex AI and how do you optimize them? **Cost**

Q3 How do you monitor Vertex AI in production? **Configuration**

Q8 How do you implement high availability with Vertex AI? **Compliance**

Q4 How do you scale Vertex AI to handle variable load? **Hardening**

Q9 How does Vertex AI integrate with a CI/CD pipeline? **Testing**

Q5 How do you secure Vertex AI in production? **SSO / IdP**

Q10 What are common pitfalls when first adopting Vertex AI? **Tuning**



ANSWER KEY & EXPLANATORY SOLUTIONS

Comprehensive verified answers, best-practice configs, security controls & reference

VERIFIED SOLUTIONS

A1 Vertex AI is a managed cloud service

Mitigation

Vertex AI is a managed cloud service that handles a specific infrastructure concern (compute, storage, messaging, AI). Using it shifts operational burden to the cloud provider while you focus on application logic.

A6 Store sensitive values in a managed secret

Observability

Store sensitive values in a managed secret store (AWS Secrets Manager, GCP Secret Manager, Azure Key Vault). Inject secrets as environment variables at runtime, never bake them into container images or source code.

A2 Access to Vertex AI is controlled via

Primitives

Access to Vertex AI is controlled via IAM roles and policies. Attach the minimum required permissions to a service account or role. Avoid long-lived access keys; use instance roles or Workload Identity Federation instead.

A7 Cost in Vertex AI typically comes from

Automation

Cost in Vertex AI typically comes from compute hours, data transfer, storage, and request counts. Right-size instances, use reserved capacity for steady-state workloads, and enable lifecycle policies to expire old data.

A3 Vertex AI emits metrics to the cloud

Hardening

Vertex AI emits metrics to the cloud provider's monitoring service (CloudWatch, Cloud Monitoring, Azure Monitor). Set alerts on key metrics (error rate, latency, capacity). Use distributed tracing to correlate across services.

A8 Deploy Vertex AI across multiple availability zones.

Governance

Deploy Vertex AI across multiple availability zones. Use managed replicas or active-active configurations. Implement health checks and automatic failover. Test resilience regularly with chaos engineering or failover drills.

A4 Vertex AI supports auto-scaling policies based on

Gotchas

Vertex AI supports auto-scaling policies based on CPU, queue depth, or custom metrics. Set minimum and maximum capacity. Horizontal scaling (more instances) is generally preferred over vertical for cloud workloads.

A9 CI/CD pipelines use the cloud CLI or

Assessment

CI/CD pipelines use the cloud CLI or SDK to deploy updates to Vertex AI after tests pass. Infrastructure-as-code (Terraform, CDK) defines Vertex AI configuration as version-controlled code deployed via the pipeline.

A5 Place Vertex AI in a private subnet

Federation

Place Vertex AI in a private subnet with no public endpoint where possible. Use VPC security groups or firewall rules to allow only required traffic. Encrypt data at rest and in transit. Rotate credentials regularly.

A10 Common mistakes include misconfigured IAM policies (too

Configuration

Common mistakes include misconfigured IAM policies (too permissive), no cost alerting, deploying only in one availability zone, and missing backups. Read the service SLA carefully to understand what the cloud provider guarantees.